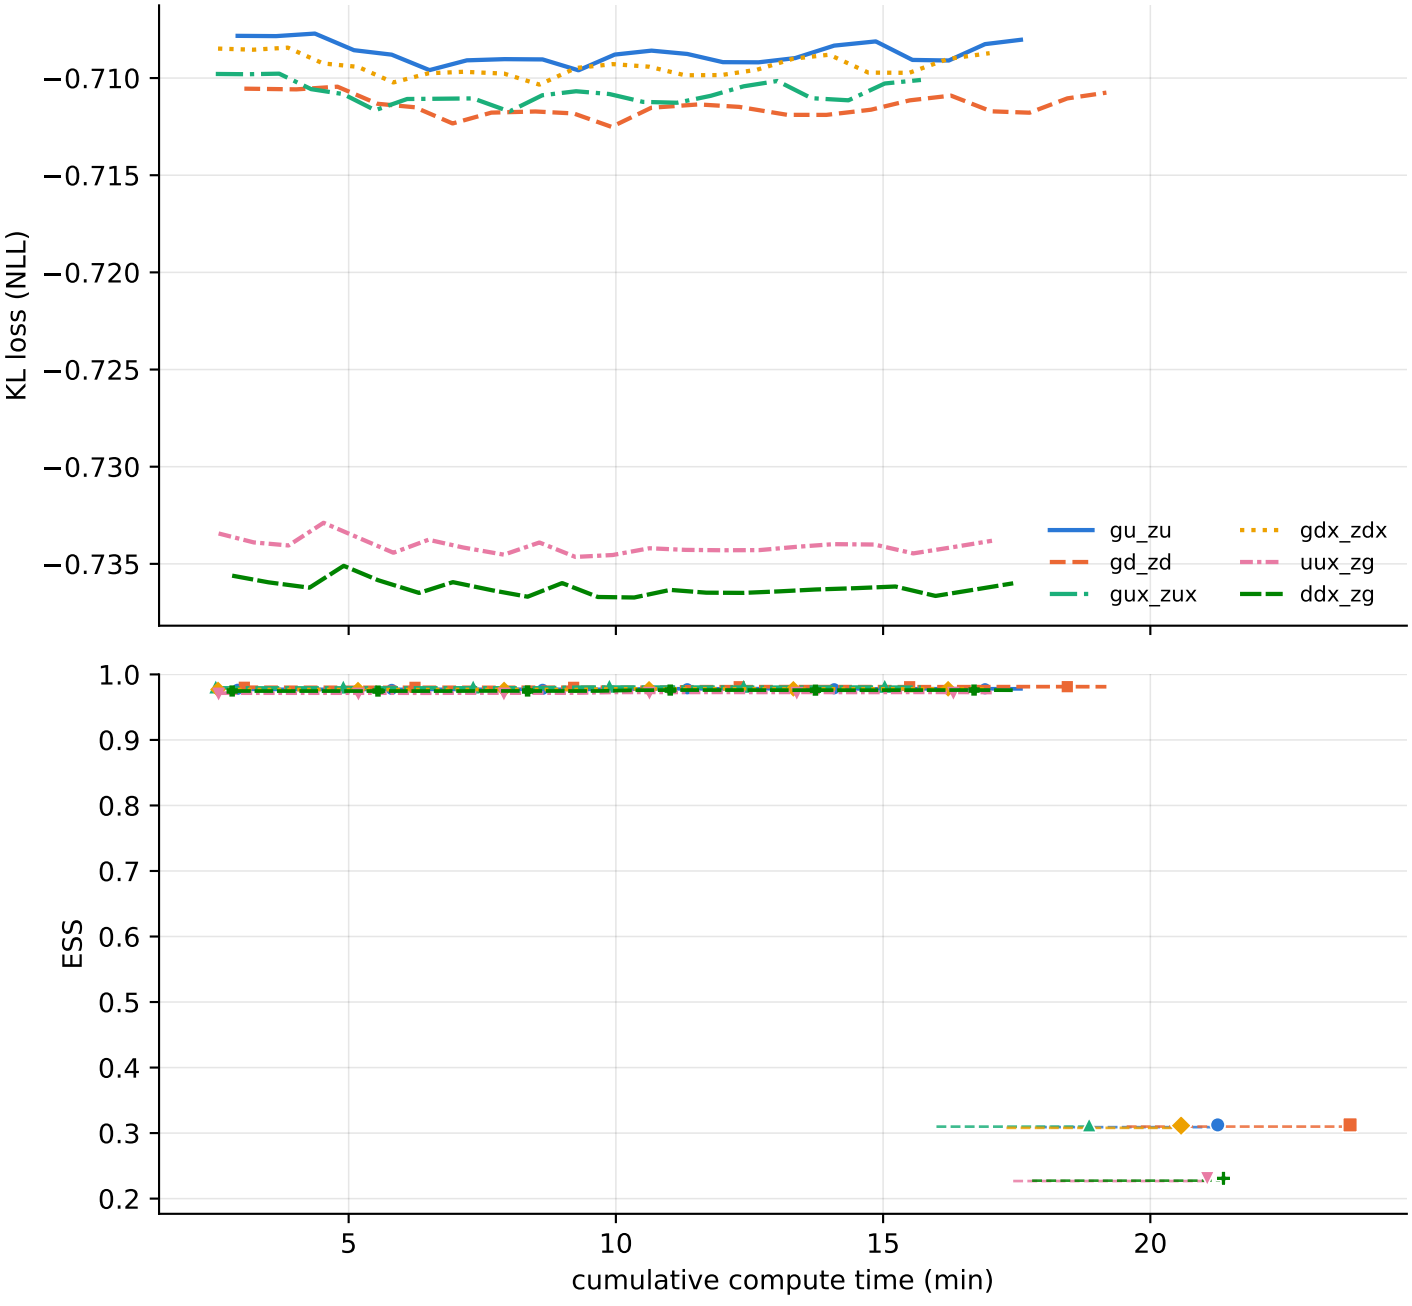
Training trajectory — high sqrt(s-hat) range
solid: KL NLL; dashed: chi2 fine-tuning ESS; dots: final ESS
KL loss (NLL)
−0.710
−0.715
−0.720
−0.725
−0.730
−0.735
gu_zu
gd_zd
gux_zux
gdx_zdx
uux_zg
ddx_zg
ESS
1.0
0.9
0.8
0.7
0.6
0.5
0.4
0.3
0.2
5
10
15
20
cumulative compute time (min)